

Random Forest Capstone

2025-10-23

Running the Code

```
library(MASS)
library(randomForest)
library(tidyverse)

data1 <- read.csv('onlyCompleteData.csv')
data1 <- data1 %>% select(WRC., k_percent, isolated_power, avg_swing_speed, squared_up_contact, avg_swing_ideal_angle_rate, exit_velocity_avg, launch_angle_avg, sweet_spot_percent, ball_hard_hit_percent, z_swing_percent, oz_swing_percent, meatball_swing_percent, pull_percent, straightaway_percent, opposite_percent, groundballs_percent, fly_linedrives_percent, popups_percent, year)

set.seed(123)
train_idx <- which(data1$year == 2024)
train <- data1[train_idx, ]
train <- train %>% select(-year)
test <- data1[-train_idx, ]
test <- test %>% select(-year)

mtry_val <- 6
ntree_values <- c(1,5,10, 25, 50, 100, 200, 500,1000)
n_reps <- 20

results <- data.frame(ntree=integer(),seed=integer(),test_MSE=double(),OOB_error=double())

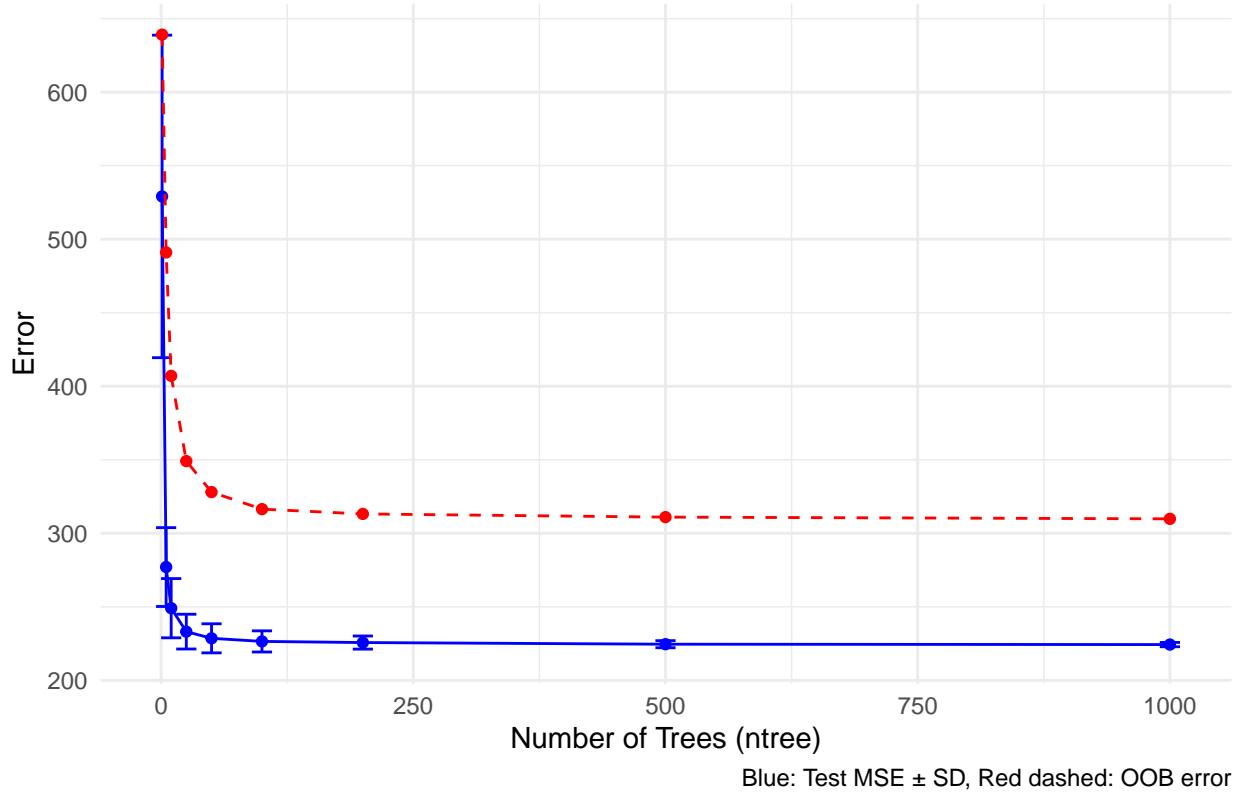
for (nt in ntree_values) { # goes through all different ntree values
  for (s in 1:n_reps) { # do many times for each ntree
    set.seed(s)
    rf_model <- randomForest(WRC. ~ ., data=train, mtry=mtry_val, ntree=nt)
    yhat.bag <- predict(rf_model, newdata=test)
    test_mse <- mean((yhat.bag - test$WRC.)^2)
    results <- rbind(results, data.frame(ntree=nt,test_MSE=test_mse,OOB_error=rf_model$mse[nt]))
  }
}

results_all <- results %>% group_by(ntree) %>% summarize(mean_test_MSE = mean(test_MSE),sd_test_MSE = sd(test_MSE),
  mean_OOB_error = mean(OOB_error),sd_OOB_error = sd(OOB_error))

ggplot(results_all, aes(x=ntree)) + geom_errorbar(aes(ymin=mean_test_MSE-sd_test_MSE, ymax=mean_test_MSE+sd_test_MSE,
  width=20, color="blue") +geom_line(aes(y=mean_test_MSE), color="blue") + geom_point(aes(y=mean_OOB_error), color="red",linetype="dashed") +
  geom_point(aes(y=mean_OOB_error), color="red") +
```

```
labs(title="Effect of ntree on Random Forest Performance",x="Number of Trees (ntree)",y="Error",caption="")
```

Effect of ntree on Random Forest Performance



```
results_all
```

```
## # A tibble: 9 x 5
##   ntree mean_test_MSE sd_test_MSE mean_OOB_error sd_OOB_error
##   <dbl>      <dbl>      <dbl>        <dbl>        <dbl>
## 1     1       529.     110.        639.        213.
## 2     5       277.     26.8       491.        67.9
## 3    10       249.     20.1       407.        33.6
## 4    25       233.     11.8       349.        25.1
## 5    50       229.      9.88      328.        19.2
## 6   100       226.      7.19      316.        11.6
## 7   200       226.      4.46      313.        8.89
## 8   500       225.      2.41      311.        5.45
## 9  1000       224.      1.47      310.        3.16

# Packages
library(MASS)
library(randomForest)
library(ggplot2)
```

```

set.seed(1)
N <- nrow(data1)
train <- which(data1$year == 2024)

mlb.train <- data1[train, ]
mlb.test <- data1[-train, ]

# Response vectors
y.test <- mlb.test$WRC.

set.seed(1)
bag.mlb <- randomForest(WRC. ~ ., data = (data1 %>% select(-year)),
                           subset = train, mtry = 12,
                           importance = TRUE)
bag.mlb

## 
## Call:
##   randomForest(formula = WRC. ~ ., data = (data1 %>% select(-year)),      mtry = 12, importance = TRUE)
##   Type of random forest: regression
##   Number of trees: 500
##   No. of variables tried at each split: 12
##
##   Mean of squared residuals: 307.8028
##   % Var explained: 36.2

yhat.bag <- predict(bag.mlb, newdata = mlb.test)

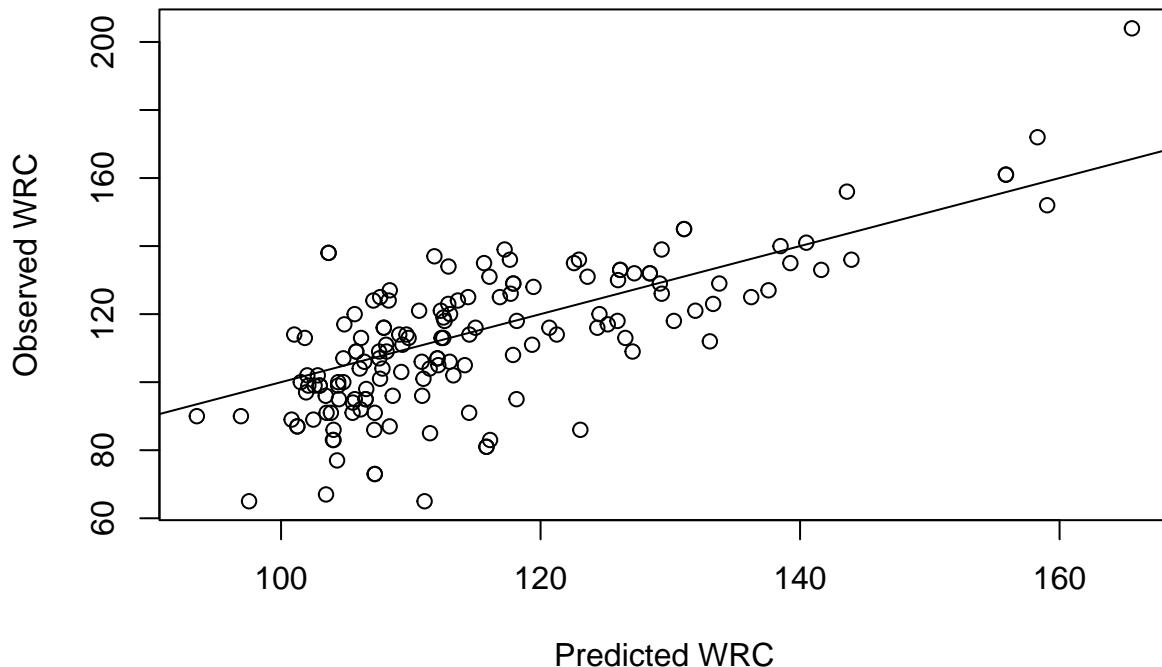
# Test MSE
mse_bag <- mean((yhat.bag - y.test)^2)
mse_bag

## [1] 217.4684

# Plot predicted vs observed
plot(yhat.bag, y.test,
      xlab = "Predicted WRC",
      ylab = "Observed WRC",
      main = "Bagging: Predicted vs Observed (Test)")
abline(0, 1)

```

Bagging: Predicted vs Observed (Test)



```
set.seed(1)
rf.mlb <- randomForest(WRC ~ ., data = (data1 %>% select(-year)),
                         subset = train, mtry = 6,
                         importance = TRUE)

yhat.rf <- predict(rf.mlb, newdata = mlb.test)
mse_rf <- mean((yhat.rf - y.test)^2)
c(mse_bag = mse_bag, mse_rf = mse_rf)

## mse_bag   mse_rf
## 217.4684 225.4536

importance(rf.mlb)

##                                     %IncMSE IncNodePurity
## k_percent                      4.0342178    1234.0679
## isolated_power                 18.9553344     8007.4562
## avg.swing_speed                6.9482535    2314.1854
## squared_up_contact              4.7277869    1246.6261
## avg.swing_length               -0.2266313     929.3655
## attack_angle                   2.4506400     935.4998
## ideal_angle_rate               -0.2542535     993.0373
## exit_velocity_avg              8.2369485    3335.9887
## launch_angle_avg                4.3392748    1190.4286
```

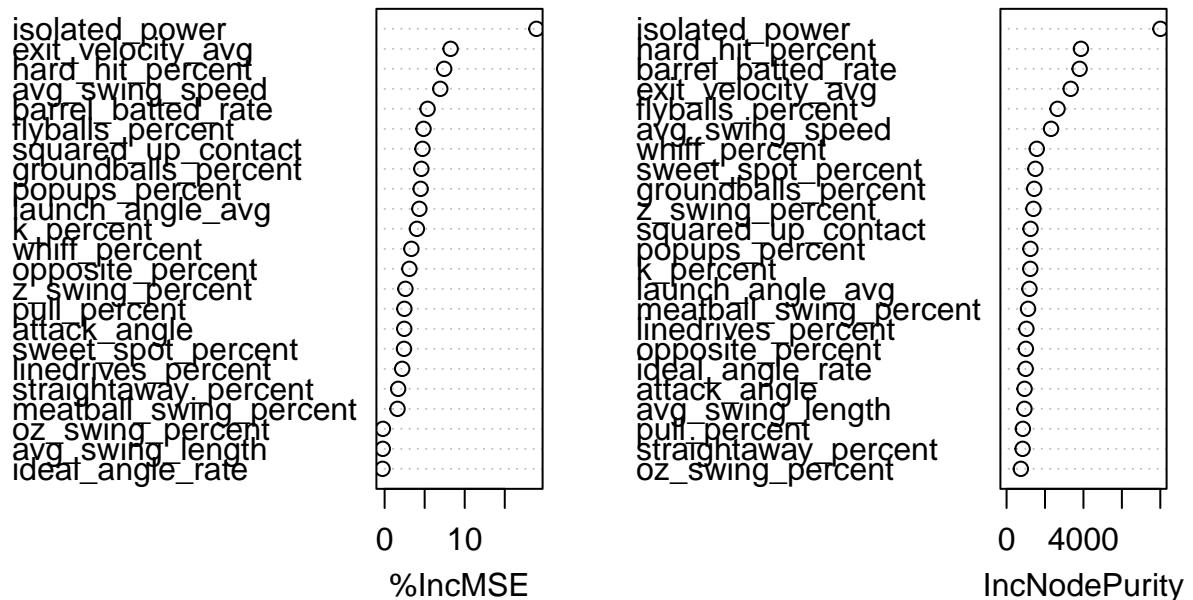
```

## sweet_spot_percent      2.4218317    1499.3724
## barrel_batted_rate     5.3643746    3800.6089
## hard_hit_percent        7.4387125    3870.8463
## z_swing_percent         2.5955117    1397.6026
## oz_swing_percent        -0.2127629   742.7318
## meatball_swing_percent  1.6135249    1117.3377
## whiff_percent           3.3409023    1571.7090
## pull_percent             2.4646820    848.0517
## straightaway_percent    1.6897713    835.0008
## opposite_percent         3.0916951    1002.2378
## groundballs_percent     4.5849651    1437.3837
## flyballs_percent         4.8595681    2653.8116
## linedrives_percent      2.1947863    1030.7874
## popups_percent           4.4976473    1240.8634

varImpPlot(rf.mlb, main = "Random Forest Variable Importance")

```

Random Forest Variable Importance



```

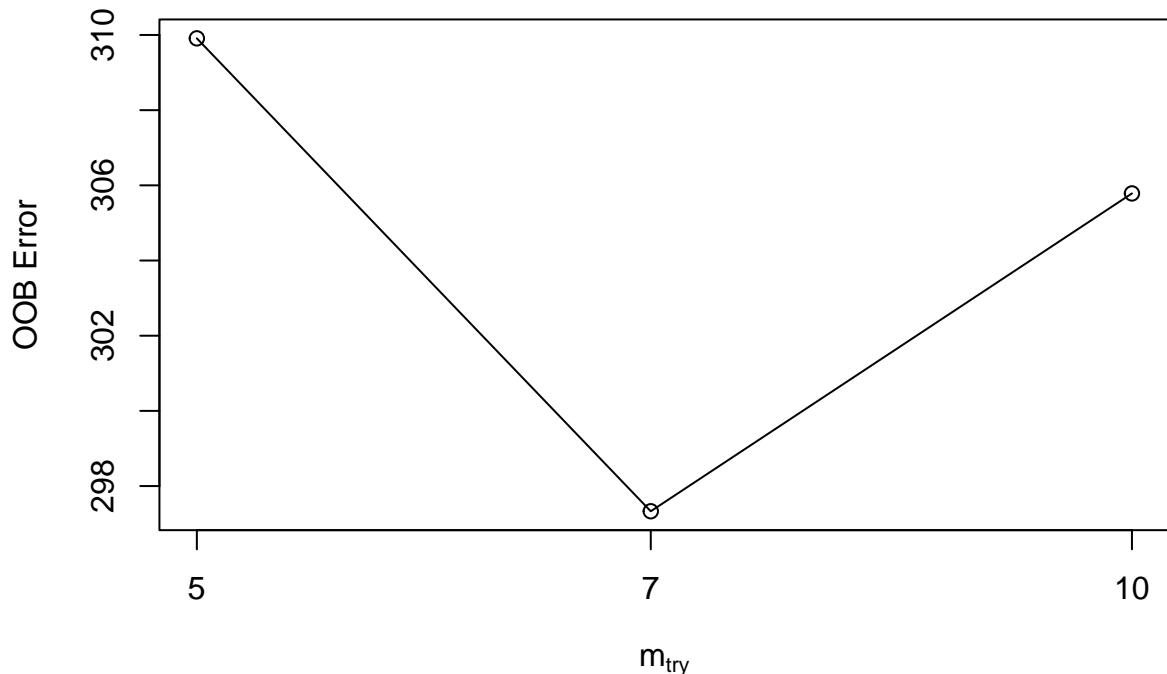
set.seed(2)
tune.out <- tuneRF(x = mlb.train[, setdiff(names(data1 %>% select(-year)), "WRC.")],
                     y = mlb.train$WRC.,
                     stepFactor = 1.5,
                     improve = 0.01,
                     ntreeTry = 500,
                     trace = TRUE,
                     plot = TRUE)

```

```

## mtry = 7  OOB error = 297.33
## Searching left ...
## mtry = 5      OOB error = 309.9098
## -0.04230923 0.01
## Searching right ...
## mtry = 10     OOB error = 305.7852
## -0.02843695 0.01

```



```

# Best mtry from tuneRF (lower OOB error is better)
best.mtry <- tune.out[which.min(tune.out[, "OOBError"]), "mtry"]
best.mtry

```

```

## [1] 7

set.seed(3)
rf.tuned <- randomForest(WRC. ~ ., data = data1 %>% select(-year),
                           subset = train, mtry = best.mtry,
                           ntree = 500, importance = TRUE)

mse_rf_tuned <- mean( (predict(rf.tuned, mlb.test) - y.test)^2 )
c(mse_bag = mse_bag, mse_rf = mse_rf, mse_rf_tuned = mse_rf_tuned)

```

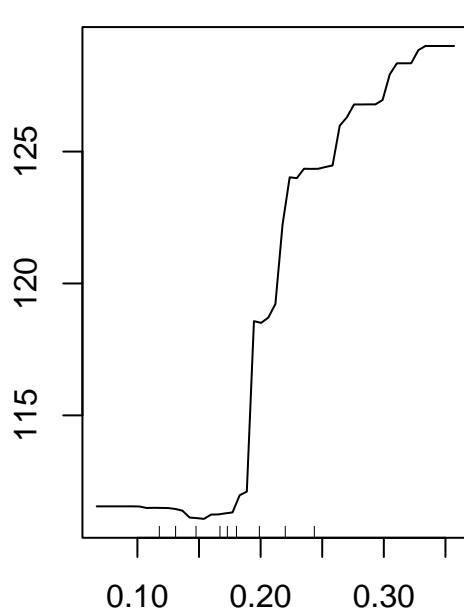
```

##      mse_bag      mse_rf mse_rf_tuned
##      217.4684    225.4536    223.0752

```

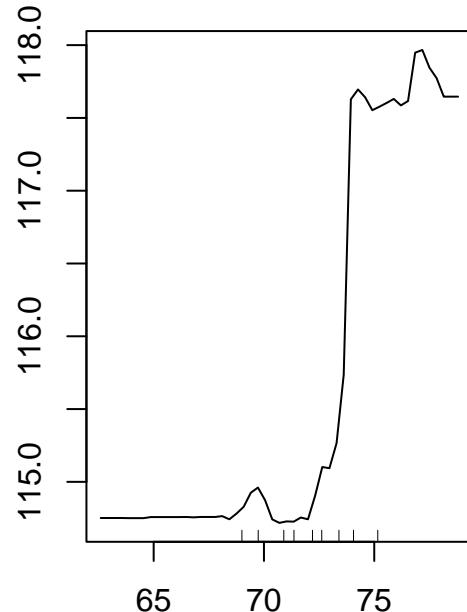
```
# Partial dependence plots give the marginal effect of a feature
par(mfrow = c(1, 2))
partialPlot(rf.tuned, pred.data = mlb.test, x.var = "isolated_power",
            main = "PDP: isolated_power")
partialPlot(rf.tuned, pred.data = mlb.test, x.var = "avg.swing_speed",
            main = "PDP: avg.swing_speed")
```

PDP: isolated_power



"isolated_power"

PDP: avg.swing_speed



"avg.swing_speed"

```
#using 2025 data

predictions <- predict(rf.tuned, mlb.test)
data1 <- read.csv('onlyCompleteData.csv')
players <- (data1 %>% filter(year==2025))$player.name
actuals <- (data1 %>% filter(year==2025))$WRC.

predict_df <- data.frame(player = players, prediction = predictions, actual = actuals)
predict_df <- predict_df %>% mutate(diff = actuals - predictions)
```