

Random Forest Capstone

2025-10-23

Lab Report

Research Question

How does the number of trees influence the performance and stability of predicting the medv of the Boston housing dataset.

Hypothesis

Initially, an increase in the number of trees will help model stability and prediction. However, there will get to a point where the increase in the number of trees will no longer improve model stability and prediction accuracy.

Parameters

The mtry will be fixed at 6. I will test ntree values of 1,5,10,25,50,100,200, and 500. Lastly, I will use test MSE and OOB error to evaluate the models.

Running the Code

```
library(MASS)
library(randomForest)
library(tidyverse)

data1 <- read.csv('onlyCompleteData.csv')

set.seed(123)
train_idx <- which(data1$year == 2024)
train <- data1[train_idx, ]
test <- data1[-train_idx, ]

mtry_val <- 6
ntree_values <- c(1,5,10, 25, 50, 100, 200, 500,1000)
n_reps <- 20

results <- data.frame(ntree=integer(),seed=integer(),test_MSE=double(),OOB_error=double())

for (nt in ntree_values) { # goes through all different ntree values
  for (s in 1:n_reps) { # do many times for each ntree
```

```


set.seed(s)
rf_model <- randomForest(WRC. ~ ., data=train, mtry=mtry_val, ntree=nt)
yhat.bag <- predict(rf_model, newdata=test)
test_mse <- mean((yhat.bag - test$WRC.)^2)
results <- rbind(results, data.frame(ntree=nt, test_MSE=test_mse, OOB_error=rf_model$mse[nt]))
}

results_all <- results %>% group_by(ntree) %>% summarize(mean_test_MSE = mean(test_MSE), sd_test_MSE = sd(test_MSE),
                                                       mean_OOB_error = mean(OOB_error), sd_OOB_error = sd(OOB_error))

ggplot(results_all, aes(x=ntree)) + geom_errorbar(aes(ymin=mean_test_MSE-sd_test_MSE, ymax=mean_test_MSE+sd_test_MSE,
                                                       width=20, color="blue") + geom_line(aes(y=mean_test_MSE), color="blue") + geom_point(aes(y=mean_test_MSE), color="blue") +
  geom_line(aes(y=mean_OOB_error), color="red", linetype="dashed") +
  geom_point(aes(y=mean_OOB_error), color="red") +
  labs(title="Effect of ntree on Random Forest Performance", x="Number of Trees (ntree)", y="Error", caption="",
       theme_minimal())

```

Effect of ntree on Random Forest Performance



```
results_all
```

```

## # A tibble: 9 x 5
##   ntree mean_test_MSE sd_test_MSE mean_OOB_error sd_OOB_error
##   <dbl>        <dbl>        <dbl>        <dbl>        <dbl>
## 1     1         467.        128.        543.        169.

```

```

## 2      5      217.     45.3      422.     101.
## 3    10      188.     28.0      366.     68.6
## 4    25      173.     12.5      288.     32.7
## 5    50      165.      8.03     269.     18.3
## 6   100      162.      4.53     259.     11.8
## 7   200      161.      3.96     256.     9.38
## 8   500      159.      2.79     253.     5.44
## 9 1000      158.      2.35     251.     3.89

```

```

# Packages
library(MASS)
library(randomForest)
library(ggplot2)

set.seed(1)
N <- nrow(data1)
train <- which(data1$year == 2024)

boston.train <- data1[train, ]
boston.test <- data1[-train, ]

# Response vectors
y.test <- boston.test$WRC.

set.seed(1)
bag.boston <- randomForest(WRC. ~ ., data = data1,
                             subset = train, mtry = 12,
                             importance = TRUE)
bag.boston

## 
## Call:
##   randomForest(formula = WRC. ~ ., data = data1, mtry = 12, importance = TRUE,      subset = train)
##   Type of random forest: regression
##   Number of trees: 500
##   No. of variables tried at each split: 12
## 
##   Mean of squared residuals: 221.0222
##   % Var explained: 54.19

yhat.bag <- predict(bag.boston, newdata = boston.test)

# Test MSE
mse_bag <- mean((yhat.bag - y.test)^2)
mse_bag

## [1] 134.0204

# Plot predicted vs observed
plot(yhat.bag, y.test,
      xlab = "Predicted WRC",


```

```

    ylab = "Observed WRC",
    main = "Bagging: Predicted vs Observed (Test)")
abline(0, 1)

```

Bagging: Predicted vs Observed (Test)



```

set.seed(1)
rf.boston <- randomForest(WRC. ~ ., data = data1,
                            subset = train, mtry = 6,
                            importance = TRUE)

```

```

yhat.rf <- predict(rf.boston, newdata = boston.test)
mse_rf <- mean((yhat.rf - y.test)^2)
c(mse_bag = mse_bag, mse_rf = mse_rf)

```

```

##  mse_bag   mse_rf
## 134.0204 159.2697

```

```
importance(rf.boston)
```

	%IncMSE	IncNodePurity
## player.name	-1.21185342	216.50609
## player_id	-0.30553609	226.55649
## year	0.00000000	0.00000
## player_age	1.74792113	135.22045
## ab	1.17875490	206.56669

## pa	1.69838655	269.41539
## k_percent	0.80992747	192.86564
## bb_percent	3.21192654	556.06730
## on_base_percent	2.89202548	516.19294
## on_base_plus_slg	5.04462262	1053.57004
## isolated_power	6.81081463	1931.56885
## b_rbi	2.82744477	834.99098
## b_walkoff	1.26048008	140.44093
## xba	-0.51729076	425.81863
## xslg	3.93312178	1000.98856
## woba	5.02477462	829.18291
## xwoba	3.88816293	791.67386
## xobp	2.30668816	577.83978
## xiso	6.02964898	1243.01092
## xbadiff	0.73517105	312.45529
## xslgdiff	0.42664003	292.91007
## wobadiff	-0.94450764	217.97177
## avg.swing.speed	2.74306040	486.04815
## fast.swing.rate	2.81913814	585.71252
## blasts.contact	3.53222158	565.80070
## blasts.swing	3.29505958	683.15174
## squared.up.contact	2.39975209	232.06789
## squared.up.swing	3.18928958	231.82940
## avg.swing.length	1.50021747	192.94863
## swords	3.44099923	288.19758
## attack.angle	1.99362675	268.20347
## attack.direction	1.34155860	202.65625
## ideal.angle.rate	2.06216931	224.76117
## vertical.swing.path	0.70847874	259.33681
## exit.velocity.avg	4.25793167	847.89712
## launch.angle.avg	1.79932686	295.53707
## sweet.spot.percent	0.85316595	316.08547
## barrel.batted.rate	1.97956923	919.99498
## solidcontact.percent	-0.88349140	164.21332
## flareburner.percent	3.99789065	450.67755
## poorlyunder.percent	2.77915379	247.40717
## poorlytopped.percent	0.23080368	383.08976
## poorlyweak.percent	-0.38193192	157.06612
## hard.hit.percent	2.07312143	793.38980
## avg.best.speed	2.41659070	820.93979
## avg.hyper.speed	3.46863482	795.82119
## z.swing.percent	1.77475360	362.66774
## z.swing.miss.percent	2.35960402	238.40913
## oz.swing.percent	-1.08675927	247.37704
## oz.swing.miss.percent	1.25835325	518.77773
## oz.contact.percent	1.71405565	398.03187
## out.zone.percent	1.42987001	139.03692
## meatball.swing.percent	1.67837611	269.61012
## meatball.percent	1.19470776	182.47170
## pitch.count.offspeed	2.23663731	838.43464
## pitch.count.fastball	-1.79706317	127.57378
## pitch.count.breaking	1.05284893	163.69261
## pitch.count	2.46452098	596.56557
## iz.contact.percent	0.54929864	216.73504


```

## in_zone_percent           1.59377562    206.50840
## edge_percent              2.06393680    213.05086
## whiff_percent             2.71800957    321.16411
## swing_percent              -0.12221427   275.69043
## pull_percent               1.02503646    249.48178
## straightaway_percent       2.72795619    308.68278
## opposite_percent            1.36699743   294.21578
## f_strike_percent           2.20202676    351.00407
## groundballs_percent        1.24962944    253.00633
## flyballs_percent            2.64526444   657.86891
## linedrives_percent         1.52491512    298.75428
## popups_percent              1.85530089   293.27783
## id                          0.30156765    253.07660
## bat_side                    1.35145563    45.05291
## side                         0.09171379   96.44633
## avg_batter_y_position      -1.59132425   196.00190
## avg_batter_x_position      2.57211059    229.47311
## avg_foot_sep                 0.98493318   219.21926
## avg_stance_angle             1.93464984   174.49810
## avg_intercept_y_vs_batter   -0.24335539   330.50249
## avg_intercept_y_vs_plate     2.21125623   395.32236
## swings_competitive          2.48255486   474.59602
## percent_swings_competitive  3.30823675   225.95232
## contact                      2.50672994   493.23295
## avg_bat_speed                4.01999102   632.26844
## hard_swing_rate                3.53081356   873.33628
## squared_up_per_bat_contact    0.85693576   317.32502
## squared_up_per_swing           2.50255060   232.48020
## blast_per_bat_contact         4.95885391   972.50777
## blast_per_swing                  5.87368452   573.52993
## swing_length                   2.28510311   338.02837
## swords_tracking                 -0.19449430  228.40300
## batter_run_value                10.72785547  3431.76524
## whiffs                        3.78217461   638.15367
## whiff_per_swing                 -1.15228190  211.63526
## batted_ball_events                1.82664782  306.43399
## batted_ball_event_per_swing     0.06043781   278.21095

varImpPlot(rf.boston, main = "Random Forest Variable Importance")


```

Random Forest Variable Importance



```
set.seed(2)
tune.out <- tuneRF(x = boston.train[, setdiff(names(data1), "WRC.")],
                     y = boston.train$WRC.,
                     stepFactor = 1.5,
                     improve = 0.01,
                     ntreeTry = 500,
                     trace = TRUE,
                     plot = TRUE)
```

```
## mtry = 32 OOB error = 195.2063
## Searching left ...
## mtry = 22 OOB error = 195.2271
## -0.0001064772 0.01
## Searching right ...
## mtry = 48 OOB error = 172.4761
## 0.1164419 0.01
## mtry = 72 OOB error = 166.5269
## 0.03449291 0.01
## mtry = 96 OOB error = 163.5571
## 0.01783359 0.01
```



```
# Best mtry from tuneRF (lower OOB error is better)
best.mtry <- tune.out[which.min(tune.out[, "OOBError"]), "mtry"]
best.mtry
```

```
## [1] 96
```


```
set.seed(3)
rf.tuned <- randomForest(WRC. ~ ., data = data1,
                           subset = train, mtry = best.mtry,
                           ntree = 500, importance = TRUE)

mse_rf_tuned <- mean( (predict(rf.tuned, boston.test) - y.test)^2 )
c(mse_bag = mse_bag, mse_rf = mse_rf, mse_rf_tuned = mse_rf_tuned)
```

```
##      mse_bag      mse_rf mse_rf_tuned
##      134.0204    159.2697    123.5992
```


```
# Partial dependence plots give the marginal effect of a feature
par(mfrow = c(1, 2))
partialPlot(rf.tuned, pred.data = boston.test, x.var = "isolated_power",
            main = "PDP: isolated_power")
partialPlot(rf.tuned, pred.data = boston.test, x.var = "avg.swing_speed",
            main = "PDP: avg.swing_speed")
```

PDP: isolated_power



"isolated_power"

PDP: avg_swing_speed



"avg_swing_speed"